

Variational Monte Carlo with Neural Quantum States

Gaussian-binary restricted Boltzmann machines for harmonic quantum dots

Marius Torsheim

Department of Physics, University of Oslo

Spring 2026

Abstract

Neural quantum states provide compact variational representations of high-dimensional wave functions. In this work I implemented a variational Monte Carlo solver where the trial wave function is a Gaussian-binary restricted Boltzmann machine. The analytical local energy, parameter-gradient estimators, and quantum force were derived in closed form and used in both brute-force Metropolis and Langevin importance sampling. The implementation was validated on non-interacting harmonic oscillator systems, for which the exact energy is $E = PD/2$ at $\omega = 1$. For one particle in one dimension, brute-force Metropolis gave $E = 0.499984 \pm 2.1 \cdot 10^{-6}$, while importance sampling gave $E = 0.499988 \pm 1.7 \cdot 10^{-6}$. For two particles in two dimensions without interaction, the corresponding estimates were $E = 1.999999 \pm 3.0 \cdot 10^{-6}$ and $E = 1.999974 \pm 3.9 \cdot 10^{-6}$. The interacting two-particle two-dimensional quantum dot was treated as the symmetric spatial ground state, corresponding to a spin-singlet state for electrons. The best completed interacting run used importance sampling, $N_h = 16$, learning rate 0.005, and gave $E = 3.0826 \pm 0.0082$, with conservative blocking error 0.0243. A larger-hidden-layer check with $N_h = 32$ gave $E = 3.0965 \pm 0.0069$ and did not improve the energy. The deviation from the exact value $E = 3.0$ is larger than the statistical uncertainty, indicating that the remaining error is dominated by optimization and ansatz limitations rather than sampling noise.

Repository: https://github.com/Torsheim/FYS4411_Project_2

I Introduction

Quantum many-body problems are difficult because the wave function depends on many coordinates and often contains non-trivial correlations. Even for few-body systems, a useful numerical method should balance physical accuracy, computational cost, and flexibility. Variational Monte Carlo (VMC) is well suited for this because it converts the ground-state problem into the optimization of a trial wave function. The variational principle guarantees that the expectation value of the Hamiltonian is an upper bound to the ground-state energy, so the quality of the trial state can be judged directly from the energy.

In this project the trial wave function is represented by a restricted Boltzmann machine (RBM). This follows the neural-network quantum state idea introduced by Carleo and Troyer [1]. The RBM is used as an explicit analytical function of the particle coordinates. Its visible nodes represent the continuous particle coordinates, while the hidden layer introduces additional nonlinear correlations. Because the coordinates are continuous, I use a Gaussian-binary RBM, with Gaussian visible nodes and binary hidden nodes.

The work has three main goals. First, the analytical expressions for the local energy, parameter derivatives, and quantum force must be derived and implemented. Second, the implementation must be validated on non-interacting harmonic oscillator systems with known exact energies. Third, the Coulomb interaction must be included for two particles in two dimensions, where an exact benchmark is available: for $\omega = 1$, the interacting ground-state energy is $E = 3.0$ a.u. [2]. The optional replacement of

the RBM by a feed-forward neural network was not attempted; instead, the focus was on obtaining a clear and tested RBM/VMC implementation.

The report is organized as a scientific report rather than as separate answers to the assignment items. Section II describes the physical systems, the RBM wave function, the analytical derivatives, and the VMC estimators. Section III explains the numerical algorithms and implementation. Section IV presents validation tests, statistical analysis, and interacting quantum-dot results. Section V summarizes the main conclusions and possible improvements. The detailed analytical derivation is included in Appendix A, while code organization and reproducibility information are summarized in Appendix B.

II Theory

A Physical systems

The Hamiltonian studied in this project is

$$\hat{H} = \sum_{p=1}^P \left(-\frac{1}{2} \nabla_p^2 + \frac{1}{2} \omega^2 r_p^2 \right) + \sum_{p < q} \frac{1}{r_{pq}}, \quad (1)$$

where P is the number of particles, D is the spatial dimension, and $r_{pq} = |\mathbf{r}_p - \mathbf{r}_q|$. Atomic units are used. The interaction term is omitted in the first numerical tests. For $\omega = 1$ and no interaction, the exact ground-state energy is

$$E_0^{(0)} = \frac{1}{2} PD. \quad (2)$$

This gives $E_0 = 0.5$ for one particle in one dimension and $E_0 = 2.0$ for two particles in two dimensions.

The main interacting benchmark is the two-particle, two-dimensional harmonic quantum dot, with

$$\mathbf{X} = (x_1, y_1, x_2, y_2), \quad (3)$$

and

$$V(\mathbf{X}) = \frac{1}{2}\omega^2(x_1^2 + y_1^2 + x_2^2 + y_2^2) + \frac{1}{\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}}. \quad (4)$$

For $\omega = 1$, the exact interacting value is $E_0 = 3.0$ a.u. [2]. This value is used as the central benchmark for the interacting calculation. Throughout the report, energies are reported in atomic units (a.u.). The oscillator frequency is fixed to $\omega = 1$ in atomic units. Monte Carlo iteration numbers, acceptance rates, numbers of hidden units, and learning rates are dimensionless. The Langevin time step Δt is the dimensionless algorithmic time step used in the importance-sampling proposal.

B Spin and exchange symmetry

The Hamiltonian in Eq. (1) is spin independent, but the interpretation of the two-particle state depends on whether the particles are treated as bosons or electrons. The RBM ansatz used here is a positive coordinate-space wave function and represents the symmetric spatial ground-state component. For two electrons in the lowest oscillator state, the corresponding total fermionic wave function must be antisymmetric under exchange. This is obtained by combining the symmetric spatial state with an antisymmetric spin singlet,

$$S = 0. \quad (5)$$

The present implementation therefore describes the spatial part of the singlet ground state. It does not include an explicit Slater determinant and does not attempt to represent triplet states or general antisymmetric many-fermion wave functions. For bosons, the same symmetric spatial ansatz can be interpreted directly as a bosonic ground-state trial function.

C Variational Monte Carlo

For a normalized trial wave function $\Psi(\mathbf{X}; \alpha)$, the variational energy is

$$E[\alpha] = \frac{\int \Psi^*(\mathbf{X}; \alpha) \hat{H} \Psi(\mathbf{X}; \alpha) d\mathbf{X}}{\int |\Psi(\mathbf{X}; \alpha)|^2 d\mathbf{X}}. \quad (6)$$

In VMC this is written as an expectation value over the probability density

$$p(\mathbf{X}) = \frac{|\Psi(\mathbf{X})|^2}{\int |\Psi(\mathbf{X})|^2 d\mathbf{X}}. \quad (7)$$

The sampled quantity is the local energy

$$E_L(\mathbf{X}) = \frac{1}{\Psi(\mathbf{X})} \hat{H} \Psi(\mathbf{X}), \quad (8)$$

so that

$$E[\alpha] = \langle E_L \rangle_p. \quad (9)$$

The energy is minimized by changing the variational parameters α of the RBM.

D Gaussian-binary RBM wave function

The visible vector $\mathbf{X} \in \mathbb{R}^M$ contains all coordinates, with

$$M = PD. \quad (10)$$

The Gaussian-binary RBM wave function used in the code is

$$\Psi(\mathbf{X}) = \frac{1}{Z} \exp \left[- \sum_{k=1}^M \frac{(X_k - a_k)^2}{2\sigma^2} \right] \prod_{j=1}^{N_h} (1 + \exp \theta_j), \quad (11)$$

with hidden pre-activations

$$\theta_j = b_j + \frac{1}{\sigma^2} \sum_{k=1}^M X_k w_{kj}. \quad (12)$$

Here $\mathbf{a}$ are visible biases, $\mathbf{b}$ are hidden biases, $\mathbf{W}$ is the weight matrix, σ is the visible Gaussian width, and N_h is the number of hidden nodes. The hidden activation probability is

$$q_j = \text{sigmoid}(\theta_j) = \frac{1}{1 + \exp(-\theta_j)}. \quad (13)$$

The logarithmic form is

$$\ln \Psi(\mathbf{X}) = \text{const.} - \sum_{k=1}^M \frac{(X_k - a_k)^2}{2\sigma^2} + \sum_{j=1}^{N_h} \ln(1 + \exp \theta_j). \quad (14)$$

The normalization constant does not enter any Metropolis ratio or coordinate derivative, and it cancels in the covariance form of the energy gradient.

E Local energy

The coordinate derivatives are most compactly written as logarithmic derivatives. Define

$$g_k = \frac{\partial \ln \Psi}{\partial X_k} = -\frac{X_k - a_k}{\sigma^2} + \frac{1}{\sigma^2} \sum_{j=1}^{N_h} w_{kj} q_j, \quad (15)$$

and

$$h_k = \frac{\partial^2 \ln \Psi}{\partial X_k^2} = -\frac{1}{\sigma^2} + \frac{1}{\sigma^4} \sum_{j=1}^{N_h} w_{kj}^2 q_j (1 - q_j). \quad (16)$$

Then

$$\frac{1}{\Psi} \frac{\partial^2 \Psi}{\partial X_k^2} = h_k + g_k^2. \quad (17)$$

The local energy is therefore

$$E_L(\mathbf{X}) = -\frac{1}{2} \sum_{k=1}^M (h_k + g_k^2) + \frac{1}{2} \omega^2 \sum_{k=1}^M X_k^2 + \sum_{p < q} \frac{1}{r_{pq}}. \quad (18)$$

For the non-interacting systems, the Coulomb term is removed. Appendix A gives the derivation in more detail.

F Parameter derivatives and optimization

The variational parameters are

$$\alpha = \{a_1, \dots, a_M, b_1, \dots, b_{N_h}, w_{11}, \dots, w_{MN_h}\}. \quad (19)$$

The logarithmic parameter derivatives are

$$O_{a_k}(\mathbf{X}) = \frac{1}{\Psi} \frac{\partial \Psi}{\partial a_k} = \frac{X_k - a_k}{\sigma^2}, \quad (20)$$

$$O_{b_j}(\mathbf{X}) = \frac{1}{\Psi} \frac{\partial \Psi}{\partial b_j} = q_j, \quad (21)$$

$$O_{w_{kj}}(\mathbf{X}) = \frac{1}{\Psi} \frac{\partial \Psi}{\partial w_{kj}} = \frac{X_k q_j}{\sigma^2}. \quad (22)$$

For any parameter α , the energy gradient is estimated as

$$\frac{\partial \langle E_L \rangle}{\partial \alpha} = 2 (\langle E_L O_\alpha \rangle - \langle E_L \rangle \langle O_\alpha \rangle). \quad (23)$$

The optimization in the production runs used ADAM [6]. In the notation of ordinary gradient descent, the update would be

$$\alpha^{(n+1)} = \alpha^{(n)} - \eta \frac{\partial \langle E_L \rangle}{\partial \alpha}, \quad (24)$$

where η is the learning rate.

G Importance sampling and quantum force

The importance-sampling version uses Langevin proposals. The quantum force is

$$F_k(\mathbf{X}) = 2 \frac{\partial \ln \Psi}{\partial X_k} = 2g_k. \quad (25)$$

A new configuration is proposed as

$$\mathbf{Y} = \mathbf{X} + D\mathbf{F}(\mathbf{X})\Delta t + \sqrt{\Delta t} \boldsymbol{\xi}, \quad (26)$$

where $D = 1/2$ and $\boldsymbol{\xi}$ is a normally distributed random vector. The Metropolis-Hastings acceptance probability includes the Green's-function ratio,

$$A(\mathbf{X} \rightarrow \mathbf{Y}) = \min \left(1, \frac{|\Psi(\mathbf{Y})|^2 G(\mathbf{Y} \rightarrow \mathbf{X})}{|\Psi(\mathbf{X})|^2 G(\mathbf{X} \rightarrow \mathbf{Y})} \right). \quad (27)$$

This is more expensive than brute-force Metropolis, but it can reduce random-walk behavior by guiding proposals toward regions where $|\Psi|^2$ is large.

H Blocking analysis

The local-energy samples are correlated because they are generated by a Markov chain. The naive standard error can therefore be too small. Blocking estimates the error by repeatedly grouping neighboring samples into blocks and computing the variance of the block averages [7]. In the present code, the sample file used for blocking contains the local-energy samples from the final optimization iteration. The first iteration of each run removes an explicit burn-in period, and later iterations continue from the previous final Markov-chain position. Thus the final sampled batch is close to equilibrium for the final optimized state, but it is not a completely independent fixed-parameter production run. This is a limitation of

the present analysis, and a stricter production protocol would freeze the optimized RBM parameters and generate a new, longer chain before blocking. In the tables below I use the largest stable or maximum blocking error as a conservative uncertainty estimate.

III Methodology and implementation

A Code structure

The implementation was written as a Python package called `nqs_vmc`. The most important files are listed in Appendix B. The central classes and functions are:

- **GaussianBinaryRBM**: evaluates $\ln \Psi$, wave-function ratios, coordinate derivatives, quantum force, and parameter derivatives.
- **local_energy.py**: evaluates Eq. (18).
- **metropolis.py**: implements brute-force Metropolis sampling.
- **importance_sampling.py**: implements the Langevin Metropolis-Hastings sampler.
- **gradients.py** and **trainer.py**: compute Eq. (23) and optimize the RBM parameters.
- **blocking.py**: computes blocking estimates for correlated samples.

The repository also contains YAML configuration files, selected result files, generated figures, generated tables, and unit tests. This organization makes it possible to reproduce runs without changing the source code.

B VMC algorithm

The following pseudocode summarizes one optimization run.

Listing 1: Pseudocode for the RBM/VMC optimization used in the project.

```
initialize RBM parameters a, b, W
for iteration = 1, ..., n_iterations:
    generate configurations X from |Psi(X)|^2
    for each sampled configuration:
        compute local energy E_L(X)
        compute logarithmic derivatives O_a,
        O_b, O_W
    estimate mean energy <E_L>
    estimate gradients 2(<E_L O> - <E_L><O>)
    update a, b, W using ADAM
    write iteration, energy, error, acceptance
    rate, gradient norm
save final parameters and local-energy samples
```

The same training loop is used for both brute-force and importance sampling; only the sampler changes.

C Tests and validation

The project was developed using tests for the RBM, derivatives, local energy, samplers, blocking routine, and reproducibility. The test suite contained 15 tests and passed before the production runs. The most important validation checks were:

- analytical derivatives were compared with finite-difference derivatives;
- the non-interacting harmonic oscillator reproduced $E = PD/2$;
- the Coulomb potential was tested for the two-particle distance;
- the same random seed reproduced the same result;
- both brute-force and importance samplers returned finite energies and reasonable acceptance rates.

The non-interacting cases are especially important because they test the RBM, local energy, sampler, optimizer, and result-writing pipeline against known analytical values.

D Run parameters

The main non-interacting runs used $\sigma = 1$, small random parameter initialization, ADAM optimization, and fixed random seeds. Brute-force Metropolis used symmetric random-walk proposals, while importance sampling used the time step Δt specified in each configuration. For the interacting system, two longer production runs were analyzed. The first used

$$\begin{aligned} P = 2, \quad D = 2, \quad \omega = 1, \\ N_h = 16, \quad \eta = 0.005, \quad \Delta t = 0.08. \end{aligned} \quad (28)$$

with 30000 samples per optimization iteration and 400 optimization iterations. Motivated by the hyperparameter sweeps, a larger-hidden-layer check used

$$\begin{aligned} P = 2, \quad D = 2, \quad \omega = 1, \\ N_h = 32, \quad \eta = 0.01, \quad \Delta t = 0.08. \end{aligned} \quad (29)$$

with 50000 samples per optimization iteration and 600 optimization iterations. This run was included to test whether the apparently favorable sweep choices combine into a better production calculation. In the interaction potential the code uses `coulomb_epsilon` as a numerical guard: if an inter-particle distance is smaller than this threshold, the local energy is set to infinity. It is not a physical softening of the Coulomb potential; for accepted configurations the interaction is still evaluated as $1/r_{12}$.

IV Results and discussion

A Overall summary

Table 1 summarizes the main completed calculations and is the numerical reference point for the rest of this section. The table makes the contrast between the essentially exact non-interacting tests and the systematically

higher interacting energies clear. The non-interacting harmonic-oscillator benchmarks are reproduced to high precision, validating the analytical local energy, parameter derivatives, samplers, and optimizer. The interacting quantum-dot calculations are less accurate, but the energies remain above the exact value, consistent with the variational principle. The larger-hidden-layer $N_h = 32$ production check did not improve the best energy estimate from the $N_h = 16$ run, illustrating that separately favorable hyperparameters do not necessarily combine into an optimal full calculation. The table separates the naive final-iteration Monte Carlo error from the more conservative blocking error whenever blocking was performed, so that statistical and systematic uncertainties are not conflated.

Table 1: Summary of the main completed energy estimates. Energies, naive errors, blocking errors, and exact values are in atomic units. A dash means that no separate blocking analysis was performed for that row.

Calculation	E	naive err.	block err.	exact
BF, $P = 1, D = 1$	0.499984	$2.1 \cdot 10^{-6}$	–	0.5
BF, $P = 2, D = 2$	1.999999	$3.0 \cdot 10^{-6}$	–	2.0
IS, $P = 1, D = 1$	0.499988	$1.7 \cdot 10^{-6}$	$1.97 \cdot 10^{-5}$	0.5
IS, $P = 2, D = 2$	1.999974	$3.9 \cdot 10^{-6}$	–	2.0
interacting short	3.1476	0.0150	–	3.0
interacting long	3.0826	0.0082	0.0243	3.0
larger N_h check	3.0965	0.0069	0.0201	3.0

B Non-interacting brute-force benchmarks

The first numerical test was the one-particle one-dimensional harmonic oscillator with two hidden nodes. The exact energy is 0.5. Figure 1 shows that the VMC energy converges rapidly to the exact value. The final value was

$$E = 0.499984 \pm 2.1 \cdot 10^{-6}. \quad (30)$$

This confirms that the analytical local energy and RBM parameter optimization are consistent for the simplest case.

figures/part_b_energy_convergence_1p1d.pdf

Figure 1: Brute-force Metropolis energy convergence for one particle in one dimension with two hidden nodes. The dashed line is the exact value $E = 0.5$.

The same non-interacting test was repeated for two particles in two dimensions. In this case $P = 2, D = 2$, and the exact energy is $E = 2.0$. The result is shown in Figure 2. The final energy was $E = 1.999999 \pm 3.0 \cdot 10^{-6}$. Table 2 gives the final numerical line of the one-dimensional benchmark and is included to document the precision and acceptance rate used for the first validation case.

figures/part_b_energy_convergence_2p2d.pdf

Figure 2: Brute-force Metropolis energy convergence for two particles in two dimensions without interaction. The exact energy is $E = 2.0$.

Table 2: Final iteration of the one-particle one-dimensional brute-force Metropolis benchmark. Energies are in atomic units.

Iter.	E [a.u.]	naive err. [a.u.]	acc.
100	0.499984	$2.1 \cdot 10^{-6}$	0.728

C Hyperparameter sensitivity of the Gaussian benchmark

The learning-rate sweep for the one-dimensional non-interacting problem is shown in Figure 3 and Table 3. Both are needed because the plotted differences are very small on the scale of the benchmark energy. All tested learning rates reproduced the exact energy well, but larger learning rates gave larger deviations and larger statistical errors. The best values in this small sweep were $\eta = 0.005$ and $\eta = 0.03$, both of which gave energies within roughly 10^{-5} of the exact result.

figures/part_b_learning_rate_sweep.pdf

Figure 3: Learning-rate sweep for the one-particle one-dimensional Gaussian benchmark. The scale is very narrow because all points are close to the exact energy.

Table 3: Learning-rate sweep for the one-dimensional non-interacting benchmark. Energies are in atomic units.

η	E [a.u.]	naive err. [a.u.]
0.005	0.499994	$1.6 \cdot 10^{-6}$
0.010	0.500050	$1.2 \cdot 10^{-5}$
0.030	0.499963	$8.5 \cdot 10^{-6}$
0.050	0.500054	$1.8 \cdot 10^{-5}$
0.080	0.500140	$4.3 \cdot 10^{-5}$

The hidden-node sweep is shown in Figure 4 and Table 4. The figure emphasizes the flat dependence on hidden-layer size, while the table gives the small numerical differences. The results show that the simple one-dimensional Gaussian benchmark does not require

many hidden nodes. One or two hidden nodes are already sufficient, which is expected because the exact ground state is a Gaussian and the Gaussian part of the RBM can represent it directly.

figures/part_b_hidden_units_sweep.pdf

Figure 4: Hidden-node sweep for the one-particle one-dimensional Gaussian benchmark. The plot uses a narrow vertical scale around $E = 0.5$.

Table 4: Hidden-node sweep for the one-dimensional non-interacting benchmark. Energies are in atomic units.

N_h	E [a.u.]	naive err. [a.u.]
1	0.500004	$5.1 \cdot 10^{-6}$
2	0.499999	$1.4 \cdot 10^{-6}$
4	0.499990	$2.0 \cdot 10^{-5}$
8	0.500031	$5.3 \cdot 10^{-5}$
16	0.500020	$4.4 \cdot 10^{-5}$

D Importance sampling benchmarks

Importance sampling was then added using the quantum force from Eq. (25). For the one-particle one-dimensional system, the final energy was

$$E = 0.499988 \pm 1.7 \cdot 10^{-6}. \quad (31)$$

The convergence curve is shown in Figure 5. The acceptance rate was very high, about 0.998, indicating that the time step was conservative.

figures/part_c_energy_convergence_1p1d_importance.pdf

Figure 5: Importance-sampling energy convergence for one particle in one dimension. The exact value is $E = 0.5$.

For the non-interacting two-particle two-dimensional case, the importance-sampling result was

$$E = 1.999974 \pm 3.9 \cdot 10^{-6}. \quad (32)$$

Figure 6 shows that this also converges to the exact value. Compared with brute-force Metropolis, the importance sampler had a much larger acceptance rate, but not necessarily a visibly smaller error in these simple systems. This is because the systems are already easy for brute-force Metropolis. The usefulness of importance sampling becomes more apparent for more complicated systems where random-walk proposals mix slowly.

figures/part_c_energy_convergence_2p2d_importance.pdf

Figure 6: Importance-sampling energy convergence for two particles in two dimensions without interaction. The exact value is $E = 2.0$.

Table 5 gives the direct numerical comparison between brute-force Metropolis and importance sampling for the two non-interacting validation systems. The two methods have comparable final energies, but their sampling behavior is different. Brute-force Metropolis with step size 1.0 gives moderate acceptance rates, especially in the four-coordinate system. Importance sampling with $\Delta t = 0.05$ gives much higher acceptance rates. For these simple Gaussian systems, however, the high acceptance rate does not lead to a dramatically smaller final error, because the brute-force sampler already performs well. This quantitative comparison supports the use of importance sampling for the interacting study, while also showing that acceptance rate alone is not a sufficient measure of estimator quality.

Table 5: Comparison of brute-force Metropolis (BF) and importance sampling (IS) on the non-interacting benchmarks. All runs used 8000 samples per optimization iteration and 1000 burn-in samples. Energies are in atomic units.

System	sampler	proposal	acc.	E [a.u.]
$P = 1, D = 1$	BF	$s = 1.0$	0.728	0.499984
$P = 1, D = 1$	IS	$\Delta t = 0.05$	0.998	0.499988
$P = 2, D = 2$	BF	$s = 1.0$	0.433	1.999999
$P = 2, D = 2$	IS	$\Delta t = 0.05$	0.993	1.999974

E Autocorrelation and blocking analysis

Blocking was applied to the local-energy samples from the non-interacting importance-sampling run. Figure 7 shows the estimated standard error as the block size increases, while Table 6 lists selected blocking levels used to choose the conservative error estimate. The conservative estimate was

$$\sigma_{\text{block}} = 1.97 \cdot 10^{-5}. \quad (33)$$

This is larger than the naive final-iteration error, as expected for correlated samples. It is therefore a more conservative uncertainty estimate.

figures/part_d_blocking_importance.pdf

Figure 7: Blocking analysis for the one-dimensional importance-sampling run. The conservative error is taken from the maximum stable block-error estimate.

Table 6: Selected blocking estimates for the non-interacting importance-sampling run.

Level	block size	error [a.u.]
5	32	$1.06 \cdot 10^{-5}$
7	128	$1.38 \cdot 10^{-5}$
9	512	$1.97 \cdot 10^{-5}$
10	1024	$1.52 \cdot 10^{-5}$

F Interacting two-particle quantum dot

The interacting two-particle two-dimensional system is the most demanding test. The exact benchmark is $E = 3.0$ a.u. The first interacting run used $N_h = 8$ and learning rate 0.01. The energy remained around 3.15 at the end of the run, as shown in Figure 8. This run is useful as a baseline, but it is not the best estimate.

figures/part_e_energy_convergence_interacting.pdf

Figure 8: Initial interacting run for two particles in two dimensions. The result remains above the exact value $E = 3.0$ and shows significant fluctuations.

The numerical values from the long interacting production runs are compared in Table 7. A longer run was then performed with $N_h = 16$, learning rate 0.005, time step $\Delta t = 0.08$, 30000 samples per iteration, and 400 optimization iterations. The convergence is shown in Figure 9. The final iteration gave

$$E = 3.08256 \pm 0.00818, \quad (34)$$

while blocking gave a conservative error of 0.0243. The energy is still above the exact value, but the improvement relative to the short run in Figure 8 is clear.

The hyperparameter sweeps suggested that larger hidden layers and a learning rate around 0.01 could improve the result. A larger-hidden-layer production check was therefore performed with $N_h = 32$, $\eta = 0.01$, $\Delta t = 0.08$, 50000 samples per iteration, and 600 optimization iterations. This run gave

$$E = 3.09649 \pm 0.00685, \quad (35)$$

with a conservative blocking error of 0.0201. This did not improve the $N_h = 16$ energy, although it produced a slightly smaller naive error and a comparable blocking error, as seen by comparing Figures 9 and 10 and the values in Table 7, where both error estimates are listed explicitly. The comparison shows that the best settings found in separate one-parameter sweeps do not necessarily

combine optimally. It also suggests that the interacting RBM optimization is limited by systematic optimization and ansatz errors, not only by Monte Carlo sampling noise.




Figure 9: Longer interacting run with $N_h = 16$. The energy decreases from about 3.23 to about 3.08, approaching but not reaching the exact value $E = 3.0$.




Figure 10: Larger-hidden-layer interacting check with $N_h = 32$, learning rate $\eta = 0.01$, and time step $\Delta t = 0.08$. The final iteration gave $E = 3.09649 \pm 0.00685$, which is higher than the best $N_h = 16$ long run despite using more hidden nodes and more samples.

Table 7: Interacting production runs. The $N_h = 16$ run gave the lowest final energy, while the $N_h = 32$ run is a larger-hidden-layer check. The naive error is the final-iteration Monte Carlo error; the blocking error is the conservative correlated-sample estimate. All energies and errors are in atomic units.

N_h	η	E	naive err.	block err.
16	0.005	3.08256	0.00818	0.0243
32	0.010	3.09649	0.00685	0.0201

The blocking analysis of the two interacting production runs is shown in Figures 11 and 12; these figures justify quoting blocking errors in addition to the naive final-iteration errors.




Figure 11: Blocking analysis for the long interacting run. The conservative blocking error is about 0.0243, larger than the naive iteration error.

figures/part_e_blocking_interacting_final.pdf

Figure 12: Blocking analysis for the larger-hidden-layer $N_h = 32$ interacting check. The conservative blocking error is about 0.0201, comparable to the long $N_h = 16$ run.

The selected numerical values from the long-run blocking curve are listed in Table 8; the largest value is the conservative error quoted for the $N_h = 16$ result.

Table 8: Selected blocking estimates for the long interacting run with $N_h = 16$.

Level	block size	error [a.u.]
5	32	0.01584
8	256	0.01561
11	2048	0.01629
13	8192	0.02427

G Interacting hyperparameter sweeps

The interacting system was also tested with several hidden-node counts, learning rates, and importance-sampling time steps. These sweeps are not full production runs, but they guide the choice of final parameters.

Figure 13 and Table 9 show the hidden-node sweep. The best value in this sweep was obtained with $N_h = 32$, which gave $E = 3.10254 \pm 0.01461$. The result suggests that the interacting system benefits from more hidden nodes than the non-interacting system, although the production comparison in Table 7 shows that a larger hidden layer alone is not sufficient to guarantee a better final optimization.

figures/part_e_hidden_units_sweep.pdf

Figure 13: Hidden-node sweep for the interacting two-particle two-dimensional system. The best tested value was $N_h = 32$.

Table 9: Hidden-node sweep for the interacting system.

N_h	E [a.u.]	naive err. [a.u.]
4	3.1881	0.0138
8	3.1859	0.0147
16	3.1878	0.0132
32	3.1025	0.0146

The learning-rate sweep is shown in Figure 14 and Table 10. These results motivated testing a larger learning rate in the final production calculation. The smallest learning rate, 0.002, converged too slowly. The best tested values were 0.01 and 0.02, both giving energies close to 3.08 in the shorter sweep. This motivated the choice of $\eta = 0.01$ for the larger-hidden-layer $N_h = 32$ run, although the combined larger-hidden-layer run did not beat the earlier $N_h = 16$ production calculation.

figures/part_e_learning_rate_sweep.pdf

Figure 14: Learning-rate sweep for the interacting system. Learning rates around 0.01 and 0.02 were best among the tested values.

Table 10: Learning-rate sweep for the interacting system.

η	E [a.u.]	naive err. [a.u.]
0.002	3.2183	0.0149
0.005	3.1709	0.0391
0.010	3.0811	0.0126
0.020	3.0830	0.0177

The time-step sweep is shown in Figure 15 and Table 11. These two objects are used together because the figure shows the trend, while the table gives the values needed to choose the production time step. Very small time steps produced high acceptance rates but poor optimization, probably because the chain moved too slowly through configuration space. The best tested time step in this sweep was $\Delta t = 0.08$, although this sweep used learning rate 0.005 and $N_h = 16$, so the conclusion should not be over-interpreted.

figures/part_e_time_step_sweep.pdf

Figure 15: Importance-sampling time-step sweep for the interacting system. The smallest time steps had very high acceptance but worse final energies.

Table 11: Time-step sweep for the interacting system.

Δt	E [a.u.]	naive err. [a.u.]
0.03	3.2771	0.0138
0.05	3.2462	0.0152
0.08	3.1648	0.0130
0.12	3.1719	0.0168

H Reliability and numerical precision

The reliability of the code is strongest for the non-interacting systems, where several independent checks agree with exact analytical results. Both brute-force and importance sampling reproduce the exact energy to roughly 10^{-5} or better. This validates the local-energy expressions, parameter derivatives, samplers, and optimization loop. The sampler comparison in Table 5 also shows that the two independent Markov-chain proposals give consistent energies, which is an additional implementation check.

The interacting system is less precise. The best long run gives an energy about 0.083 above the exact value, or roughly 2.8% relative error. The larger-hidden-layer $N_h = 32$ check gives an energy about 0.0965 above the exact value, or roughly 3.2% relative error. These deviations are larger than the statistical error estimates, meaning that the dominant error is not sampling noise alone. It is most likely a combination of incomplete optimization, limited expressive power of the present RBM ansatz, and hyperparameter sensitivity. The energies remain above the exact value, which is consistent with the variational principle, but the interacting result is not yet fully optimized.

The blocking curves in Figures 11 and 12, together with Table 8, show that the interacting statistical uncertainty is much larger than in the non-interacting validation case. The blocking samples were taken from the final sampled batch of each optimization run after the burn-in procedure described in Section III. This is sufficient for a conservative project-level error estimate, but a stricter final calculation should freeze the optimized parameters and generate an independent, longer production chain before blocking. The acceptance rates in importance sampling were around 0.99 for many runs. Such high acceptance rates are safe, but they can also indicate that the time step is conservative. The time-step sweep supports this interpretation: the smallest time steps had the highest acceptance rates but did not give the best energies. Future production runs should therefore tune the time step and learning rate together rather than independently and repeat the best settings with several random seeds.

V Conclusion

A Gaussian-binary restricted Boltzmann machine was implemented as a neural quantum state for variational Monte Carlo calculations. The analytical expressions for the local energy, quantum force, and parameter gradients were derived and implemented in closed form. The implementation was validated against non-interacting harmonic oscillator benchmarks, where the exact energy is $E = PD/2$. The final non-interacting estimates were within about 10^{-5} of the exact answers for both brute-force Metropolis and importance sampling. The two-particle electron calculation was interpreted as the symmetric spatial part of a spin-singlet state, so no explicit antisymmetric spatial factor was included.

For the interacting two-particle two-dimensional quantum dot, the exact benchmark is $E = 3.0$ a.u. The best completed long run gave $E = 3.0826 \pm 0.0082$, with a

conservative blocking error of 0.0243. A larger-hidden-layer run with $N_h = 32$, $\eta = 0.01$, and more samples gave $E = 3.0965 \pm 0.0069$, with conservative blocking error 0.0201, and therefore did not improve the energy. The interacting calculations are clear improvements over the initial run, but they are still not fully converged to the exact result. Because the difference from the exact value is several times larger than the blocking error, the remaining discrepancy should be interpreted as a systematic optimization or ansatz error, not simply a Monte Carlo error.

The most concrete improvements would be to perform a joint sweep over $(N_h, \eta, \Delta t)$ with repeated random seeds, introduce a decaying learning-rate schedule, freeze the optimized parameters and run a separate fixed-parameter production chain for the final blocking analysis, and replace ordinary ADAM optimization by stochastic reconfiguration or a natural-gradient method. A more explicitly correlated trial state, for example by multiplying the RBM by a Jastrow factor or by implementing the optional neural-network/Jastrow ansatz, would also be a natural next step. These changes directly target the dominant limitations observed here: optimization sensitivity and insufficient correlation in the interacting state.

A Analytical derivative details

The central identity used in the local energy is

$$\frac{1}{\Psi} \frac{\partial^2 \Psi}{\partial X_k^2} = \frac{\partial^2 \ln \Psi}{\partial X_k^2} + \left(\frac{\partial \ln \Psi}{\partial X_k} \right)^2. \quad (36)$$

For the Gaussian-binary RBM, Eq. (14) gives

$$\frac{\partial \ln \Psi}{\partial X_k} = -\frac{X_k - a_k}{\sigma^2} + \frac{1}{\sigma^2} \sum_j w_{kj} q_j, \quad (37)$$

where $q_j = \text{sigmoid}(\theta_j)$. Since

$$\frac{\partial q_j}{\partial X_k} = q_j(1 - q_j) \frac{w_{kj}}{\sigma^2}, \quad (38)$$

the second logarithmic derivative is

$$\frac{\partial^2 \ln \Psi}{\partial X_k^2} = -\frac{1}{\sigma^2} + \frac{1}{\sigma^4} \sum_j w_{kj}^2 q_j(1 - q_j). \quad (39)$$

Substitution into the kinetic-energy operator gives Eq. (18). The same logarithmic derivative g_k gives the quantum force $F_k = 2g_k$ used in the Langevin proposal.

B Code and reproducibility

The repository is organized with a source package, scripts, tests, configuration files, and report files. The project can be installed with

```
python -m venv .venv
source .venv/bin/activate
pip install -e .
```

The tests can be run with

```
pytest -q
```

The main scripts used for the results are

```
python scripts/run_experiment.py --config
  configs/part_b_1p1d_2h_bruteforce.yaml
python scripts/run_experiment.py --config
  configs/part_c_1p1d_2h_importance.yaml
python scripts/run_experiment.py --config
  configs/part_e_2p2d_interacting.yaml
```

The generated histories, blocking samples, figures, and tables are stored under the **results/** and **report/** folders.

Appendix: Use of AI/LLM Tools

ChatGPT was used as a limited supporting tool during this project. The tool was used during parts of the project work in 2026 through the ChatGPT web interface. The implementation, numerical experiments, figures, tables, and final scientific interpretation are my own responsibility.

For the code, ChatGPT was used mainly for support while developing and cleaning the project. This included help with discussing repository organization, interpreting terminal errors, checking Git and L^AT_EX commands, improving file names and structure, and making comments and code organization easier to read. I also used ChatGPT for debugging support and for suggestions about useful consistency checks. All code suggestions were reviewed before use, and the submitted code was tested by running the project workflow and unit tests.

For the written report, ChatGPT was used as a language and feedback tool. It helped with proofreading, improving wording, checking notation consistency, and giving feedback on whether the report addressed the course evaluation criteria. I reviewed and edited the suggestions before including them in the final text.

For text assistance, level 0 means no LLM assistance, level 1 means language-only assistance, level 2 means editorial assistance, and level 3 means generative drafting. In this report the assistance was mainly level 1, with some level 2 editorial feedback for the AI declaration and general report structure. No report section was included without my own review and approval.

For code assistance, level 0 means no LLM assistance, level 1 means debugging or explanation assistance, level 2 means short snippets or small blocks, level 3 means skeleton generation, and level 4 means substantial file-level generation. In this project the assistance was mainly level 1–2: debugging support, command checking, code cleanup suggestions, comments, naming, and short code suggestions. ChatGPT was not used to fabricate numerical results, choose final data, or replace my own calculations and scientific interpretation.

All reported values were generated by running the project code in the GitHub repository. The numerical results were checked against exact diagonalization, analytic limits, unit tests, or other relevant consistency checks. I take responsibility for the submitted code, figures, tables, numerical values, and scientific interpretation, and I can

explain the equations, algorithms, and results used in the final submission.

References

- [1] G. Carleo and M. Troyer, “Solving the quantum many-body problem with artificial neural networks,” *Science* **355**, 602–606 (2017).
- [2] M. Taut, “Two electrons in an external oscillator potential: Particular analytic solutions of a Coulomb correlation problem,” *Physical Review A* **48**, 3561–3566 (1993).
- [3] G. Torlai and R. G. Melko, “Machine-learning quantum states in the NISQ era,” *Annual Review of Condensed Matter Physics* **11**, 325–344 (2020).
- [4] G. E. Hinton, “A practical guide to training restricted Boltzmann machines,” Technical Report UTMML TR 2010-003, University of Toronto (2010).
- [5] H. Saito, “Method to solve quantum few-body problems with artificial neural networks,” *Journal of the Physical Society of Japan* **87**, 074002 (2018).
- [6] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” International Conference on Learning Representations (2015).
- [7] H. Flyvbjerg and H. G. Petersen, “Error estimates on averages of correlated data,” *Journal of Chemical Physics* **91**, 461–466 (1989).
- [8] Computational Physics II FYS4411/FYS9411, “Project 2: Machine learning for quantum many-body problems,” University of Oslo, Spring 2026. <https://github.com/CompPhysics/ComputationalPhysics2/tree/gh-pages/doc/Projects/2026/Project2/Project2ML>
- [9] Computational Physics II FYS4411/FYS9411, “Lecture notes on Boltzmann machines,” University of Oslo. https://compphysics.github.io/ComputationalPhysics2/doc/LectureNotes/_build/html/boltzmannmachines.html